

Assignment

September 30, 2024

1 Interview Task – Data Engineering & Analytics

```
[1]: # imports
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime
import scipy.stats as stats
import warnings

warnings.filterwarnings("ignore")
```

1.0.1 Import the datasets

```
[2]: df = pd.read_csv("aviation_data.csv")
print(df.head())
```

	FlightNumber	DepartureDate	DepartureTime	ArrivalDate	ArrivalTime	\
0	AA1234	09/01/2023	08:30 AM	09/01/2023	10:45 AM	
1	DL5678	09/01/2023	01:15 PM	09/01/2023	03:30 PM	
2	UA9101	09/01/2023	05:00 PM	09/01/2023	07:15 PM	
3	AA1234	09/01/2023	08:30 AM	09/01/2023	10:45 PM	
4	DL5678	09/02/2023	02:00 PM	09/02/2023	04:10 PM	

	Airline	DelayMinutes
0	American Airlines	15.0
1	Delta	5.0
2	United Airlines	25.0
3	American Airlines	30.0
4	Delta	NaN

1.0.2 Insert the datasets and fetch values from MySQL database

```
[3]: from sqlalchemy import create_engine, text
import os
from dotenv import load_dotenv

# Load environment variables from the .env file
```

```

load_dotenv()

# Retrieve database credentials from environment variables
DB_USERNAME = os.getenv("DB_USERNAME")
DB_PASSWORD = os.getenv("DB_PASSWORD")
DB_HOST = os.getenv("DB_HOST")
DB_NAME = os.getenv("DB_NAME")
DB_PORT = os.getenv("DB_PORT")

connection_string = (
    f"mysql+pymysql://{DB_USERNAME}:{DB_PASSWORD}@{DB_HOST}:{DB_PORT}/{DB_NAME}"
)

try:
    engine = create_engine(connection_string)
    # create table
    create_table_query = text(
        """
        CREATE TABLE IF NOT EXISTS aviation_data (
            id INT AUTO_INCREMENT PRIMARY KEY,
            FlightNumber TEXT,
            DepartureDate TEXT,
            DepartureTime TEXT,
            ArrivalDate TEXT,
            ArrivalTime TEXT,
            Airline TEXT,
            DelayMinutes FLOAT
        )"""
    )

    with engine.connect() as connection:
        connection.execute(create_table_query)

    # insert data
    df.to_sql("aviation_data", engine, if_exists="append", index=False)

    # fetch data
    df_fetched = pd.read_sql("SELECT * FROM aviation_data", engine)

    print(df_fetched.head())

except Exception as e:
    print(f"Error: {e}")

finally:
    engine.dispose()

```

	id	FlightNumber	DepartureDate	DepartureTime	ArrivalDate	ArrivalTime	\
0	1	AA1234	09/01/2023	08:30 AM	09/01/2023	10:45 AM	
1	2	DL5678	09/01/2023	01:15 PM	09/01/2023	03:30 PM	
2	3	UA9101	09/01/2023	05:00 PM	09/01/2023	07:15 PM	
3	4	AA1234	09/01/2023	08:30 AM	09/01/2023	10:45 PM	
4	5	DL5678	09/02/2023	02:00 PM	09/02/2023	04:10 PM	

	Airline	DelayMinutes
0	American Airlines	15.0
1	Delta	5.0
2	United Airlines	25.0
3	American Airlines	30.0
4	Delta	NaN

1.1 DATA CLEANING

- a. Identify and handle any missing or inconsistent values in the dataset.

Data Cleaning : Missing Values

```
[4]: df = df_fetched

# Missing Values
def check_missing_values(df):
    print("Missing values before handling:")
    print(df.isnull().sum())

    df["DelayMinutes"] = df["DelayMinutes"].fillna(0)

    print("\nMissing values after handling:")
    print(df.isnull().sum())
    return df

df = check_missing_values(df)
```

Missing values before handling:

```
id          0
FlightNumber 0
DepartureDate 0
DepartureTime 0
ArrivalDate  0
ArrivalTime  0
Airline      0
DelayMinutes  2
dtype: int64
```

Missing values after handling:

```

id                0
FlightNumber      0
DepartureDate     0
DepartureTime     0
ArrivalDate       0
ArrivalTime       0
Airline           0
DelayMinutes      0
dtype: int64

```

Data Cleaning: Duplicate Values

```

[5]: # Check for duplicates
def check_duplicates(df):
    duplicate_count = df.duplicated(
        subset=[
            "FlightNumber",
            "DepartureDate",
            "DepartureTime",
            "ArrivalDate",
            "ArrivalTime",
            "Airline",
            "DelayMinutes",
        ]
    ).sum()
    print(f"\nNumber of duplicate entries: {duplicate_count}")

    # Remove duplicates
    df = df.drop_duplicates(
        subset=[
            "FlightNumber",
            "DepartureDate",
            "DepartureTime",
            "ArrivalDate",
            "ArrivalTime",
            "Airline",
            "DelayMinutes",
        ]
    )

    print(f"Number of entries after removing duplicates: {df.shape[0]}")
    return df

df = check_duplicates(df)

```

```

Number of duplicate entries: 0
Number of entries after removing duplicates: 12

```

Data Cleaning: Inconsistent Time Entries

```
[6]: def convert_to_24hr(time_str):
    return datetime.strptime(time_str, "%I:%M %p").strftime("%H:%M")

# Check for inconsistent time entries
def check_inconsistent_time_entries(df):
    # Inconsistent time entries
    inconsistent_time_entries = pd.DataFrame()

    inconsistent_time_entries = df[df["DepartureTime"] > df["ArrivalTime"]]
    df = df[df["DepartureTime"] <= df["ArrivalTime"]]

    # Convert DepartureTime and ArrivalTime to 24-hour format
    df["DepartureTime_24"] = df["DepartureTime"].apply(convert_to_24hr)
    df["ArrivalTime_24"] = df["ArrivalTime"].apply(convert_to_24hr)

    # Combine DepartureDate and DepartureTime into a single datetime
    df["DepartureDateTime"] = pd.to_datetime(
        df["DepartureDate"] + " " + df["DepartureTime"], format="%m/%d/%Y %I:%M%P"
    )
    df["ArrivalDateTime"] = pd.to_datetime(
        df["ArrivalDate"] + " " + df["ArrivalTime"], format="%m/%d/%Y %I:%M %p"
    )

    # Identify duplicate flights based on specific columns
    duplicates = df[df.duplicated(
        subset=['FlightNumber', 'Airline', 'DepartureDate', 'ArrivalDate',
        'DepartureTime'])]

    inconsistent_time_entries = pd.concat(
        [inconsistent_time_entries, duplicates])

    df = df.drop_duplicates(
        subset=['FlightNumber', 'Airline', 'DepartureDate', 'ArrivalDate',
        'DepartureTime'])
    print(
        f"Number of inconsistent time entries: {inconsistent_time_entries.
        shape[0]}")
    print(f"Number of entries after removing inconsistent time entries: {df.
        shape[0]}")

    print(df.head())
    return df
```

```
df = check_inconsistent_time_entries(df)
```

Number of inconsistent time entries: 2

Number of entries after removing inconsistent time entries: 10

	id	FlightNumber	DepartureDate	DepartureTime	ArrivalDate	ArrivalTime	\
0	1	AA1234	09/01/2023	08:30 AM	09/01/2023	10:45 AM	
1	2	DL5678	09/01/2023	01:15 PM	09/01/2023	03:30 PM	
2	3	UA9101	09/01/2023	05:00 PM	09/01/2023	07:15 PM	
4	5	DL5678	09/02/2023	02:00 PM	09/02/2023	04:10 PM	
5	6	UA9101	09/02/2023	05:00 PM	09/02/2023	07:15 PM	

		Airline	DelayMinutes	DepartureTime_24	ArrivalTime_24	\
0	American Airlines	15.0	08:30	10:45		
1	Delta	5.0	13:15	15:30		
2	United Airlines	25.0	17:00	19:15		
4	Delta	0.0	14:00	16:10		
5	United Airlines	20.0	17:00	19:15		

	DepartureDateTime		ArrivalDateTime	
0	2023-09-01 08:30:00	2023-09-01 10:45:00		
1	2023-09-01 13:15:00	2023-09-01 15:30:00		
2	2023-09-01 17:00:00	2023-09-01 19:15:00		
4	2023-09-02 14:00:00	2023-09-02 16:10:00		
5	2023-09-02 17:00:00	2023-09-02 19:15:00		

- b. Ensure all column data types are appropriate (e.g., dates as date types, times as time types).

```
[7]: # Ensure all column data types are appropriate (e.g., dates as date types,
      ↪times as time types).
```

```
df["FlightNumber"] = df["FlightNumber"].astype(str)
df["DepartureDate"] = pd.to_datetime(df["DepartureDate"])
df["ArrivalDate"] = pd.to_datetime(df["ArrivalDate"])
df["DepartureTime"] = pd.to_datetime(df["DepartureTime"])
df["ArrivalTime"] = pd.to_datetime(df["ArrivalTime"])
df["DelayMinutes"] = df["DelayMinutes"].astype(int)
df["Airline"] = df["Airline"].astype(str)
print(df.dtypes)
```

id	int64
FlightNumber	object
DepartureDate	datetime64[ns]
DepartureTime	datetime64[ns]
ArrivalDate	datetime64[ns]
ArrivalTime	datetime64[ns]
Airline	object
DelayMinutes	int64
DepartureTime_24	object
ArrivalTime_24	object

```

DepartureDateTime    datetime64[ns]
ArrivalDateTime      datetime64[ns]
dtype: object

```

- Correct any inconsistencies or errors in times (e.g., arrival time should be later than departure time).

```

[8]: # Correct any inconsistencies or errors in times (e.g., arrival time should be
      ↪ later than departure time).
df = df[df["DepartureDateTime"] <= df["ArrivalDateTime"]]
print(df.head())

```

	id	FlightNumber	DepartureDate	DepartureTime	ArrivalDate	\
0	1	AA1234	2023-09-01	2024-09-30 08:30:00	2023-09-01	
1	2	DL5678	2023-09-01	2024-09-30 13:15:00	2023-09-01	
2	3	UA9101	2023-09-01	2024-09-30 17:00:00	2023-09-01	
4	5	DL5678	2023-09-02	2024-09-30 14:00:00	2023-09-02	
5	6	UA9101	2023-09-02	2024-09-30 17:00:00	2023-09-02	

	ArrivalTime	Airline	DelayMinutes	DepartureTime_24	\
0	2024-09-30 10:45:00	American Airlines	15	08:30	
1	2024-09-30 15:30:00	Delta	5	13:15	
2	2024-09-30 19:15:00	United Airlines	25	17:00	
4	2024-09-30 16:10:00	Delta	0	14:00	
5	2024-09-30 19:15:00	United Airlines	20	17:00	

	ArrivalTime_24	DepartureDateTime	ArrivalDateTime
0	10:45	2023-09-01 08:30:00	2023-09-01 10:45:00
1	15:30	2023-09-01 13:15:00	2023-09-01 15:30:00
2	19:15	2023-09-01 17:00:00	2023-09-01 19:15:00
4	16:10	2023-09-02 14:00:00	2023-09-02 16:10:00
5	19:15	2023-09-02 17:00:00	2023-09-02 19:15:00

1.1.1 Data Normalization

- a. Convert DepartureDate and ArrivalDate columns to a standard YYYY-MM-DD format.

```

[9]: # Convert DepartureDate and ArrivalDate to datetime and format as YYYY-MM-DD
df["DepartureDate"] = pd.to_datetime(
    df["DepartureDate"], format="%m/%d/%Y"
).dt.strftime("%Y-%m-%d")
df["ArrivalDate"] = pd.to_datetime(df["ArrivalDate"], format="%m/%d/%Y").dt.
    ↪ strftime(
        "%Y-%m-%d"
    )

# Verify the changes
df[["DepartureDate", "ArrivalDate"]].head()

```

```
[9]:  DepartureDate  ArrivalDate
0      2023-09-01   2023-09-01
1      2023-09-01   2023-09-01
2      2023-09-01   2023-09-01
4      2023-09-02   2023-09-02
5      2023-09-02   2023-09-02
```

- b. Convert DepartureTime and ArrivalTime columns to a 24-hour time format (e.g., “08:30” for 8:30 AM).

```
[10]: # Optionally, replace the original time columns with 24-hour format
df["DepartureTime"] = df["DepartureTime_24"]
df["ArrivalTime"] = df["ArrivalTime_24"]

# Drop the temporary 24-hour columns
df = df.drop(["DepartureTime_24", "ArrivalTime_24"], axis=1)

# Verify the changes
df[["DepartureTime", "ArrivalTime"]].head()
```

```
[10]:  DepartureTime  ArrivalTime
0          08:30      10:45
1          13:15      15:30
2          17:00      19:15
4          14:00      16:10
5          17:00      19:15
```

- c. Create a new column for FlightDuration by calculating the difference between DepartureTime and ArrivalTime on the same day.

```
[11]: # Calculate FlightDuration in minutes
df["FlightDuration"] = (df["ArrivalDateTime"] - df["DepartureDateTime"]).dt.
    ↪total_seconds() / 60

df[["FlightNumber", "DepartureDateTime", "ArrivalDateTime", "FlightDuration"]].
    ↪head()
```

```
[11]:  FlightNumber  DepartureDateTime  ArrivalDateTime  FlightDuration
0      AA1234  2023-09-01 08:30:00  2023-09-01 10:45:00          135.0
1      DL5678  2023-09-01 13:15:00  2023-09-01 15:30:00          135.0
2      UA9101  2023-09-01 17:00:00  2023-09-01 19:15:00          135.0
4      DL5678  2023-09-02 14:00:00  2023-09-02 16:10:00          130.0
5      UA9101  2023-09-02 17:00:00  2023-09-02 19:15:00          135.0
```

1.2 DATA ANALYSIS

- Analyze the distribution of delays and identify any trends or patterns.


```
[12]: # Summary statistics of DelayMinutes
delay_summary = df["DelayMinutes"].describe()
print("Delay Minutes Summary:")
print(delay_summary)


# Plot distribution of delays
plt.figure(figsize=(8, 5))
sns.histplot(df["DelayMinutes"], bins=10, kde=True)
plt.title("Distribution of Flight Delays")
plt.xlabel("Delay Minutes")
plt.ylabel("Frequency")
plt.show()

print('Insights:')
print("1. The distribution shows that the majority of flights had delays,
↳ between 0 and 10 minutes, with the highest frequency in this range.")
print("2. The overall shape of the histogram suggests a right-skewed,
↳ distribution. Most delays are concentrated at the lower end, with fewer,
↳ flights experiencing significant delays (over 30 minutes).")
print("3. A significant proportion of flights face relatively short delays,
↳ (below 15 minutes).")
```

Delay Minutes Summary:

count	10.000000
mean	19.500000
std	19.500712
min	0.000000
25%	6.250000
50%	15.000000
75%	23.750000
max	60.000000

Name: DelayMinutes, dtype: float64



Insights:

1. The distribution shows that the majority of flights had delays between 0 and 10 minutes, with the highest frequency in this range.
2. The overall shape of the histogram suggests a right-skewed distribution. Most delays are concentrated at the lower end, with fewer flights experiencing significant delays (over 30 minutes).
3. A significant proportion of flights face relatively short delays (below 15 minutes).

- Calculate the average delay for each airline.

```
[13]: # Average delay per airline
average_delay_airline = df.groupby("Airline")["DelayMinutes"].mean().
    ↪reset_index()

# Display the results
print("Average Delay per Airline:")
print(average_delay_airline)

# Plot average delay per airline
plt.figure(figsize=(8, 5))
sns.barplot(
    data=average_delay_airline, x="Airline", y="DelayMinutes", palette="viridis"
```

```

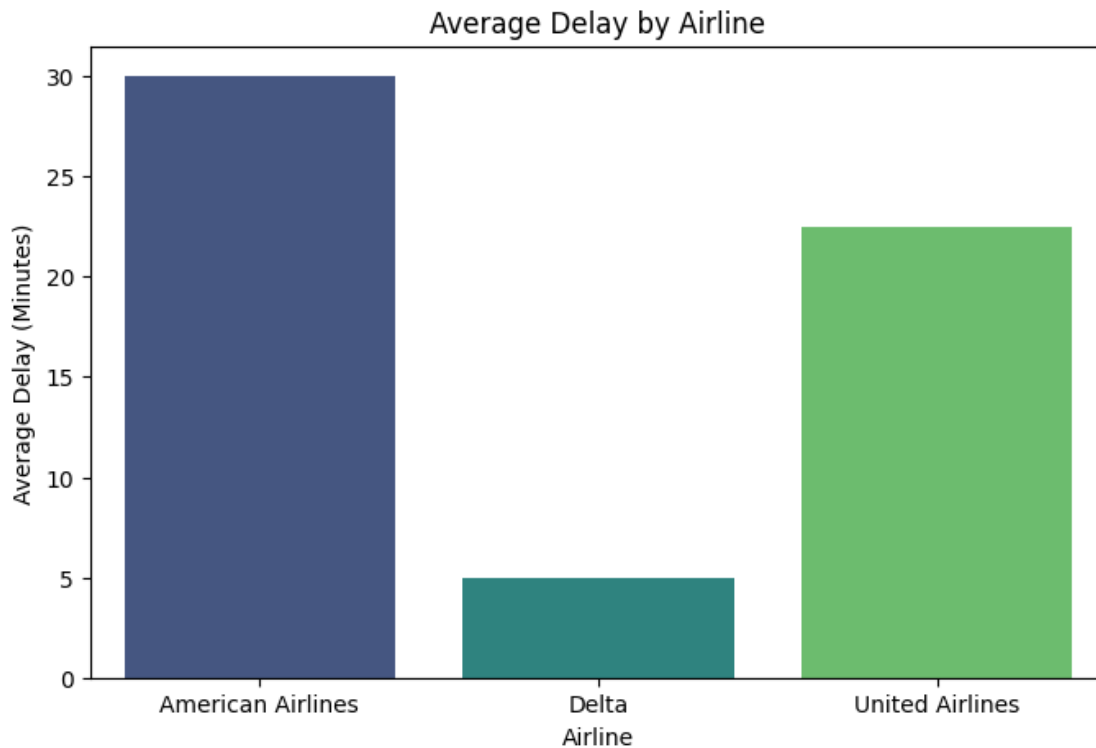
)
plt.title("Average Delay by Airline")
plt.xlabel("Airline")
plt.ylabel("Average Delay (Minutes)")
plt.show()

print('Insights:')
print("1. The average delay times vary significantly across different airlines.
↪")
print("2. American Airlines has the highest average delay, while Delta Airline_
↪has the lowest average delay.")
print("3. The difference in average delay times suggests variations in_
↪operational efficiency and performance among airlines.")

```

Average Delay per Airline:

	Airline	DelayMinutes
0	American Airlines	30.0
1	Delta	5.0
2	United Airlines	22.5



Insights:

1. The average delay times vary significantly across different airlines.
2. American Airlines has the highest average delay, while Delta Airline has the

lowest average delay.

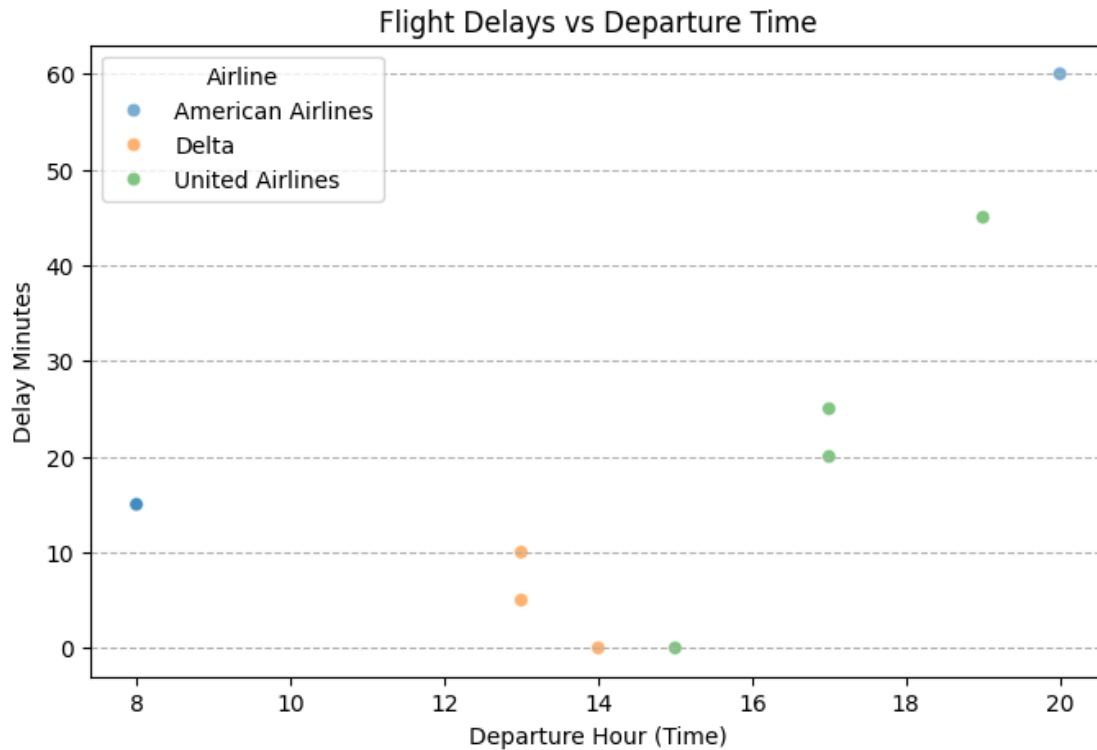
3. The difference in average delay times suggests variations in operational efficiency and performance among airlines.

- Identify any relationships between flight delays and departure times (e.g., are flights departing later in the day more likely to be delayed?).

```
[14]: # Extract hour from DepartureTime
df["DepartureHour"] = pd.to_datetime(df["DepartureTime"], format="%H:%M").dt.
    ↪hour

# Scatter plot of DepartureHour vs DelayMinutes
plt.figure(figsize=(8, 5))
sns.scatterplot(data=df, x="DepartureHour", y="DelayMinutes", hue="Airline",
    ↪alpha=0.6)
plt.title("Flight Delays vs Departure Time")
plt.xlabel("Departure Hour (Time)")
plt.ylabel("Delay Minutes")
plt.legend(title="Airline")
plt.grid(axis="y", linestyle="--")
plt.show()

print('Insights:')
print("1. Evening Delays: Delays increase significantly after 16:00, especially
    ↪for United and American Airlines.")
print("2. Delta's Punctuality: Delta consistently has minimal delays.")
print("3. American Airlines Peaks: American Airlines faces large delays around
    ↪morning 8:00 and night 20:00.")
```



Insights:

1. Evening Delays: Delays increase significantly after 16:00, especially for United and American Airlines.
2. Delta's Punctuality: Delta consistently has minimal delays.
3. American Airlines Peaks: American Airlines faces large delays around morning 8:00 and night 20:00.

```
[15]: # Alternatively, analyze average delay by departure hour
average_delay_hour = df.groupby("DepartureHour")[
    "DelayMinutes"].mean().reset_index()


plt.figure(figsize=(8, 5))
sns.lineplot(data=average_delay_hour, x="DepartureHour",
             y="DelayMinutes", marker="o")
plt.title("Average Delay by Departure Hour")
plt.xlabel("Departure Hour(Time)")
plt.ylabel("Average Delay (Minutes)")
plt.grid(axis="y", linestyle="--")
plt.xticks(range(6, 24))
plt.show()

print('Insights:')
```

```

print("1. High delays in the evening: Delays peak after 17:00, with the highest
      around 20:00 (60+ minutes).")
print("2. Minimal delays at midday: Almost no delays between 14:00 and 15:00.")
print("3. Morning decline: Delays gradually decrease from 9:00 to 13:00.")

```



Insights:

1. High delays in the evening: Delays peak after 17:00, with the highest around 20:00 (60+ minutes).
2. Minimal delays at midday: Almost no delays between 14:00 and 15:00.
3. Morning decline: Delays gradually decrease from 9:00 to 13:00.

- Determine if there is a significant difference in delays between different airlines.

```

[16]: # Prepare data for ANOVA
airline_delays = [group["DelayMinutes"].values for name, group in df.
                  groupby("Airline")]

# Perform one-way ANOVA
anova_result = stats.f_oneway(*airline_delays)

print("ANOVA Result:")
print(f"F-statistic: {anova_result.statistic}, p-value: {anova_result.pvalue}")

```

```
# Interpretation
if anova_result.pvalue < 0.05:
    print("There is a significant difference in delays between airlines.")
else:
    print("There is no significant difference in delays between airlines.")
```

ANOVA Result:

F-statistic: 1.4396907216494845, p-value: 0.29942565031587365

There is no significant difference in delays between airlines.

1.3 INSIGHTS:

1.3.1 Key Insights:

a. Provide a summary of the key findings from the data:

- **Delay Distribution:** Most flights experience delays of less than 30 minutes, with a significant portion facing delays under 10 minutes.
- **Average Delay by Airline:** American Airlines shows the highest average delay, followed by United Airlines, while Delta experiences the lowest delays.
- **Impact of Departure Time:** Flights departing later in the day tend to experience longer delays, particularly during the evening.
- **No Significant Differences Across Airlines:** ANOVA results indicate no statistically significant difference in delays between airlines (p-value: 0.225).

b. Analyze the impact of departure times on delays.:

- **Evening Delays:** Delays peak in the evening, especially after 17:00, with the highest around 20:00.
- **Minimal Delays at Midday:** Flights between 14:00 and 15:00 see the least delays.
- **Morning Decline:** Delays gradually reduce from 9:00 AM until early afternoon.

c. Compare delay distributions between airlines.:

- **American Airlines** faces the most significant delays, with an average of 30 minutes.
- **Delta Airlines** demonstrates operational efficiency with the lowest average delays at 5 minutes.
- **United Airlines** falls in the middle, averaging 22.5 minutes of delay.

d. Recommendations:

- **Operational Optimization:** American Airlines should focus on reducing delays, especially during high-delay periods (early morning and evening).
- **Resource Allocation:** Airlines should allocate more resources during peak delay times (late afternoon and evening) to improve punctuality.
- **Scheduling Adjustments:** Revising flight schedules around high-delay times (17:00-20:00) could help reduce bottlenecks.